
Beyond Parallel Corpora

Philipp Koehn

29 October 2020



data and machine learning

Supervised and Unsupervised



- We framed machine translation as a supervised machine learning task
 - training examples with labels
 - here: input sentences with translation
 - structured prediction: output has to be constructed in several steps■
- Unsupervised learning
 - training examples without labels
 - here: just sentences in the input language
 - we will also look at using just sentences output language■
- Semi-supervised learning
 - some labeled training data
 - some unlabeled training data (usually more)■
- Self-training
 - make predictions on unlabeled training data
 - use predicted labeled as supervised translation data

Transfer Learning



- Learning from data similar to our task
- Other language pairs
 - first, train a model on different language pair
 - then, train on the targeted language pair
 - or: train jointly on both
- Multi-Task training
 - train on a related task first
 - e.g., part-of-speech tagging
- Share some or all of the components

using monolingual data

Using Monolingual Data



- Language model
 - trained on large amounts of target language data
 - better fluency of output
- Key to success of statistical machine translation
- Neural machine translation
 - integrate neural language model into model
 - create artificial data with backtranslation

Adding a Language Model

- Train a separate language model
- Add as conditioning context to the decoder
- Recall state progression in the decoder
 - decoder state s_i
 - embedding of previous output word Ey_{i-1}
 - input context c_i

$$s_i = f(s_{i-1}, Ey_{i-1}, c_i)$$

- Add hidden state of neural language model s_i^{LM}

$$s_i = f(s_{i-1}, Ey_{i-1}, c_i, s_i^{\text{LM}})$$

- Pre-train language model
- Leave its parameters fixed during translation model training

- Balance impact of language model vs. translation model

- Learn a scaling factor (gate) $\text{gate}_i^{\text{LM}} = f(s_i^{\text{LM}})$

- Use it to scale values of language model state


$$\bar{s}_i^{\text{LM}} = \text{gate}_i^{\text{LM}} \times s_i^{\text{LM}}$$

- Use this scaled language model state for decoder state

$$s_i = f(s_{i-1}, Ey_{i-1}, c_i, \bar{s}_i^{\text{LM}})$$

Back Translation


- Monolingual data is parallel data that misses its other half
- Let's synthesize that half



- Steps
 1. train a system in reverse language translation
 2. use this system to translate target side monolingual data
→ synthetic parallel corpus
 3. combine generated synthetic parallel data with real parallel data to build the final system
- Roughly equal amounts of synthetic and real data
- Useful method of domain adaptation

Iterative Back Translation

- Quality of backtranslation system matters
- Build a better backtranslation system ... with backtranslation



Iterative Back Translation

- Example


German–English	Back	Final
no back-translation	-	29.6
*10k iterations	10.6	29.6 (+0.0)
*100k iterations	21.0	31.1 (+1.5)
convergence	23.7	32.5 (+2.9)
re-back-translation	27.9	33.6 (+4.0)

* = limited training of back-translation system

Round Trip Training

- We could iterate through steps of
 - train system
 - create synthetic corpus
- Dual learning: train models in both directions together
 - translation models $F \rightarrow E$ and $E \rightarrow F$
 - take sentence \mathbf{f}
 - translate into sentence \mathbf{e}'
 - translate that back into sentence \mathbf{f}'
 - training objective: \mathbf{f} should match \mathbf{f}'
- Setup could be fooled by just copying ($\mathbf{e}' = \mathbf{f}$)
 - \Rightarrow score \mathbf{e}' with a language for language E
 - add language model score as cost to training objective


Round Trip Training



- Copy Target
 - if no good neural machine translation system to start with
 - just copy target language text to the source
- Forward Translation
 - synthesize training data in same direction as training
 - self-training (inferior but sometimes successful)


unsupervised machine translation

Monolingual Embedding Spaces



- Embedding spaces for different languages have similar shape
- Intuition: relationship between *dog*, *cat*, and *lion*, independent of language
- How can we rotate the triangle to match up?

Matching Embedding Spaces



- Seed lexicon of identically spelled words, numbers, names
- Adversarial training method: discriminator predicts [Conneau et al., 2018]
- Match matrices with word similarity scores: Vecmap [Artetxe et al., 2018]

- Translation model
 - induced word translations (nearest neighbors of mapped embeddings)
 - statistical phrase translation table (probability \simeq similarity)
 - Language model
 - target side monolingual data
 - estimate statistical n-gram language model
- ⇒ Statistical phrase-based machine translation system

Synthetic Training Data

- Create synthetic parallel corpus
 - monolingual text in source language
 - translate with inferred system: translations in target language
- Recall: EM algorithm
 - predict data: generate translation for monolingual corpus
 - predict model: estimate model from synthetic data
 - iterate this process, alternate between language directions
- Increasingly use neural machine translation model to synthesize data

multiple language pairs

Multiple Language Pairs

- There are more than two languages in the world
- We may want to build systems for many language pairs
- Typical: train separate models for each
- Joint training

Multiple Input Languages

- Example
 - German–English
 - French–English
- Concatenate training data
- Joint model benefits from exposure to more English data
- Shown beneficial in low resource conditions
- Do input languages have to be related? (maybe not)

Multiple Output Languages

- Example
 - French–English
 - French–Spanish
- Concatenate training data
- Given a French input sentence, how specify output language? ■
- Indicate output language with special tag

[ENGLISH] *N'y a-t-il pas ici deux poids, deux mesures?*


⇒ *Is this not a case of double standards?*

[SPANISH] *N'y a-t-il pas ici deux poids, deux mesures?*

⇒ *¿No puede verse con toda claridad que estamos utilizando un doble rasero?*

Zero Shot Translation

- Example
 - German–English
 - French–English
 - French–Spanish
- We want to translate
 - German–Spanish



- Train on
 - German–English
 - French–English
 - French–Spanish
- Specify translation

[SPANISH] *Messen wir hier nicht mit zweierlei Maß?*

⇒ *¿No puede verse con toda claridad que estamos utilizando un doble rasero?*

Algorithms

Google's AI just created its own universal 'language'

The technology used in Google Translate can identify hidden material between languages to create what's known as interlingua

By **MATT BURGESS**

23 Nov 2016

Table 5: Portuguese→Spanish BLEU scores using various models.

	Model	Zero-shot	BLEU
(a)	PBMT bridged	no	28.99
(b)	NMT bridged	no	30.91
(c)	NMT Pt→Es	no	31.50
(d)	Model 1 (Pt→En, En→Es)	yes	21.62
(e)	Model 2 (En↔{Es, Pt})	yes	24.75
(f)	Model 2 + incremental training	no	31.77

- Bridged: pivot translation Portuguese → English → Spanish
- Model 1 and 2: Zero shot training
- Model 2 + incremental training: use of some training data in language pair


Sharing Components

- So far: generic neural machine translation model
- Maybe better: separate systems with shared components
 - encoder shared in models with same input language.
 - decoder shared in models with same output language.
 - attention mechanism shared in all models
- Sharing = same parameters, updates from any language pair training
- No need to mark output language

Massively Multilingual Training

- Scaling up multilingual machine translation for more languages
 - many-to-English
 - English-to-many
 - many-to-many
- Mainly motivated by improving low-resource language pairs
- Move towards larger models


Translation Quality for 103 Languages



(source: Google)

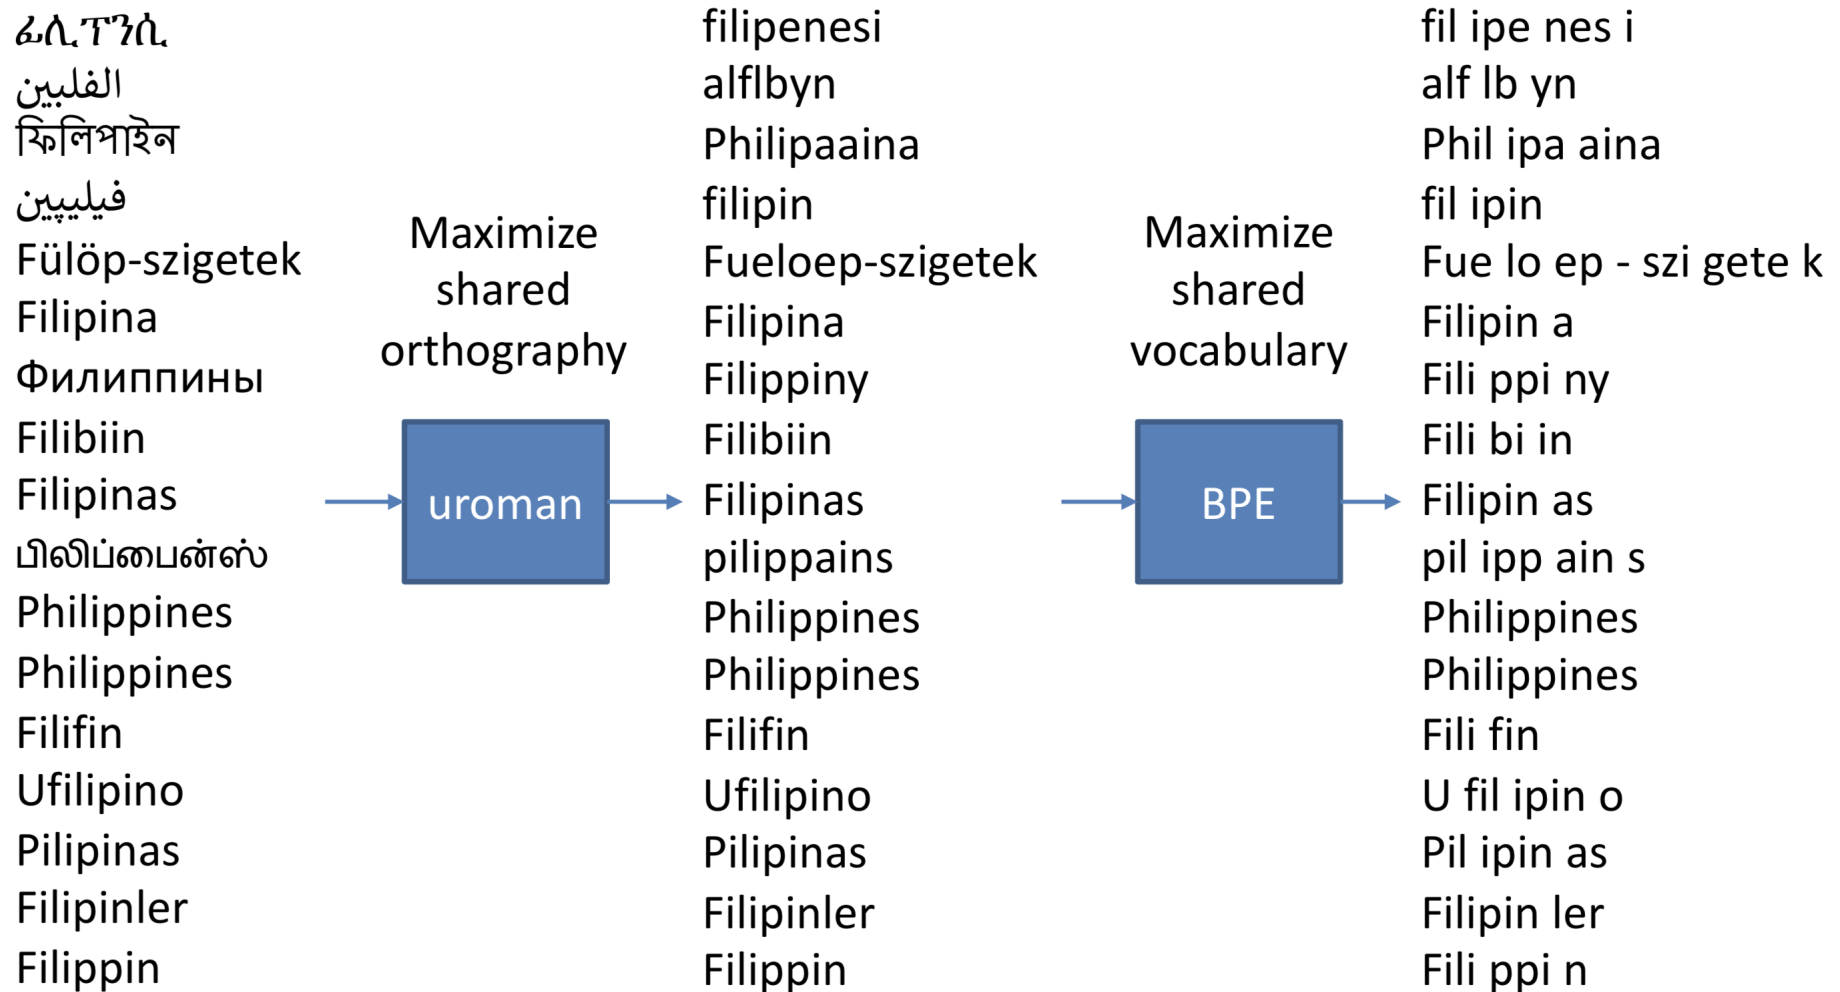
Gains with Multilingual Training

31



(source: Google)

Romanization



(source: USC/ISI)

Facebook

Introducing the First AI Model That Translates 100 Languages Without Relying on English

October 19, 2020

By Angela Fan, Research Assistant

- 7.5 billion sentences for 100 languages (mined from web-crawled data)
- Model with 15 billion parameters
- Improvements especially for low resource languages

multi-task training

- Our translation models: generic sequence-to-sequence models
- Same model used for many other tasks
 - sentiment detection
 - grammar correction
 - semantic inference
 - summarization
 - question answering
 - speech recognition
- For all these tasks, we need to learn basic properties of language
 - word embeddings
 - contextualize word representations in encoder
 - language model aspects of decoder
- Why re-invent the wheel each time?

Training on Related Tasks

- Train model on several tasks
- Maybe shared and task-specific components
- System learns general facts about language
 - informed by many different tasks
 - useful for many different tasks

Pre-Training Word Embeddings

- Let us keep it simple...
 - Neural machine translation models use word embeddings
 - encoding of input words
 - encoding of output words
 - Word embeddings can be trained on vast amounts of monolingual data
- ⇒ pre-train word embeddings and initialize model with them
- Not very successful so far
 - monolingual word embeddings trained on language model objectives
 - for machine translation, different similarity aspects may matter more
 - e.g., *teacher* and *teaching* similar in MT, not in LM

- Pre-training other components of the translation model
- Decoder
 - language model, informed by input context
 - pre-train as language model on monolingual data
 - input context vector set to zero
- Encoder
 - also structures like a language model
(however, not optimized to predict following words)
 - pre-train as language model on monolingual data

Monolingual Pre-Training

- Initial training of neural machine translation model on monolingual data
- Replace some input word sequences with <pad> (30% of words)
- Train model MASKED → TEXT on both source and target text
- Reorder sentences (each training example has 3 sentences)

<en> Advanced NLP techniques master class " how <pad> " </s>

3rd <pad> : 18 </s>

Results <pad> 40 of 729



3rd grade : 18 </s>

Advanced NLP techniques master class " how to with clients " </s>

Results 1 – 40 of 729

Multi-Task Training

- Multiple end-to-end tasks that share common aspects
 - need to encode an input word sequence
 - produce an output word sequence■
- May have very different input/output
 - sentiment detection: output is sentiment value
 - part-of-speech tagging: output is tag sequence
 - syntactic parsing: output is recursive parse structure (may be linearized)
 - semantic parsing: output is logical form, database query, or AMR
 - grammar correction: input is error-prone text
 - question answering: needs to be also informed by knowledge base
 - speech recognition: input is sequence of acoustic features■
- Input and output in the same language, may be mostly copied
 - grammar correction, automatic post-editing
 - question answering, semantic inference

- Train a single model for all tasks
- Positive results with joint training of
 - part-of-speech tagging
 - named entity recognition
 - syntactic parsing
 - semantic analysis.
- Tasks may share just some components