

---

# IBM Model 1 and the EM Algorithm

Philipp Koehn

4 September 2025



- How to translate a word → look up in dictionary

**Haus** — house, building, home, household, shell.

- Multiple translations
  - some more frequent than others
  - for instance: **house**, and **building** most common
  - special cases: **Haus** of a **snail** is its **shell**
- Note: In all lectures, we translate from a foreign language into English

# Collect Statistics



Look at a parallel corpus (German text along with English translation)

| Translation of <i>Haus</i> | Count |
|----------------------------|-------|
| house                      | 8,000 |
| building                   | 1,600 |
| home                       | 200   |
| household                  | 150   |
| shell                      | 50    |

# Estimate Translation Probabilities



3


Maximum likelihood estimation

$$p_f(e) = \begin{cases} 0.8 & \text{if } e = \text{house}, \\ 0.16 & \text{if } e = \text{building}, \\ 0.02 & \text{if } e = \text{home}, \\ 0.015 & \text{if } e = \text{household}, \\ 0.005 & \text{if } e = \text{shell}. \end{cases}$$

# Alignment



- In a parallel text (or when we translate), we align words in one language with the words in the other



- Word positions are numbered 1–4

# Alignment Function




- Formalizing alignment with an alignment function
- Mapping an English target word at position  $i$  to a German source word at position  $j$  with a function  $a : i \rightarrow j$
- Example

$$a : \{1 \rightarrow 1, 2 \rightarrow 2, 3 \rightarrow 3, 4 \rightarrow 4\}$$

# Reordering

Words may be reordered during translation



$$a : \{1 \rightarrow 3, 2 \rightarrow 4, 3 \rightarrow 2, 4 \rightarrow 1\}$$

# One-to-Many Translation




A source word may translate into multiple target words



$$a : \{1 \rightarrow 1, 2 \rightarrow 2, 3 \rightarrow 3, 4 \rightarrow 4, 5 \rightarrow 4\}$$

# Dropping Words


Words may be dropped when translated  
(German article **das** is dropped)



$$a : \{1 \rightarrow 2, 2 \rightarrow 3, 3 \rightarrow 4\}$$

# Inserting Words

- Words may be added during translation
  - The English **just** does not have an equivalent in German
  - We still need to map it to something: special NULL token



$$a : \{1 \rightarrow 1, 2 \rightarrow 2, 3 \rightarrow 3, 4 \rightarrow 0, 5 \rightarrow 4\}$$

- Generative model: break up translation process into smaller steps
  - IBM Model 1 only uses lexical translation
- Translation probability
  - for a foreign sentence  $\mathbf{f} = (f_1, \dots, f_{l_f})$  of length  $l_f$
  - to an English sentence  $\mathbf{e} = (e_1, \dots, e_{l_e})$  of length  $l_e$
  - with an alignment of each English word  $e_j$  to a foreign word  $f_i$  according to the alignment function  $a : j \rightarrow i$

$$p(\mathbf{e}, a | \mathbf{f}) = \frac{\epsilon}{(l_f + 1)^{l_e}} \prod_{j=1}^{l_e} t(e_j | f_{a(j)})$$

- parameter  $\epsilon$  is a normalization constant

# Example

das

| $e$   | $t(e f)$ |
|-------|----------|
| the   | 0.7      |
| that  | 0.15     |
| which | 0.075    |
| who   | 0.05     |
| this  | 0.025    |

Haus

| $e$       | $t(e f)$ |
|-----------|----------|
| house     | 0.8      |
| building  | 0.16     |
| home      | 0.02     |
| household | 0.015    |
| shell     | 0.005    |

ist

| $e$    | $t(e f)$ |
|--------|----------|
| is     | 0.8      |
| 's     | 0.16     |
| exists | 0.02     |
| has    | 0.015    |
| are    | 0.005    |

klein

| $e$    | $t(e f)$ |
|--------|----------|
| small  | 0.4      |
| little | 0.4      |
| short  | 0.1      |
| minor  | 0.06     |
| petty  | 0.04     |

$$\begin{aligned} p(e, a|f) &= \frac{\epsilon}{5^4} \times t(\text{the}|\text{das}) \times t(\text{house}|\text{Haus}) \times t(\text{is}|\text{ist}) \times t(\text{small}|\text{klein}) \\ &= \frac{\epsilon}{5^4} \times 0.7 \times 0.8 \times 0.8 \times 0.4 \\ &= 0.00029\epsilon \end{aligned}$$



# finding translations

# Centauri-Arcturan Parallel Text

1a. ok-voon ororok sprok .  
1b. at-voon bichat dat .

---

2a. ok-drubel ok-voon anak plok sprok .  
2b. at-drubel at-voon pippat rrat dat .

---

3a. erok sprok izok hihok ghirok .  
3b. totat dat arrat vat hilat .

---

4a. ok-voon anak drok brok jok .  
4b. at-voon krat pippat sat lat .

---

5a. wiwok farok izok stok .  
5b. totat jjat quat cat .

---

6a. lalok sprok izok jok stok .  
6b. wat dat krat quat cat .

---

7a. lalok farok ororok lalok sprok izok enemok .  
7b. wat jjat bichat wat dat vat eneat .

---

8a. lalok brok anak plok nok .  
8b. iat lat pippat rrat nnat .

---

9a. wiwok nok izok kantok ok-yurp .  
9b. totat nnat quat oloat at-yurp .

---

10a. lalok mok nok yorok ghirok klok .  
10b. wat nnat gat mat bat hilat .

---

11a. lalok nok crrrok hihok yorok zanzanok .  
11b. wat nnat arrat mat zanzanat .

---

12a. lalok rarok nok izok hihok mok .  
12b. wat nnat forat arrat vat gat .

---

Translation challenge: **farok crrrok hihok yorok klok kantok ok-yurp**

(from Knight (1997): Automating Knowledge Acquisition for Machine Translation)



# em algorithm




- We would like to estimate the lexical translation probabilities  $t(e|f)$  from a parallel corpus
- ... but we do not have the alignments
- Chicken and egg problem
  - if we had the *alignments*,  
→ we could estimate the *parameters* of our generative model
  - if we had the *parameters*,  
→ we could estimate the *alignments*

# EM Algorithm

- Incomplete data
  - if we had *complete data*, would could estimate *model*
  - if we had *model*, we could fill in the *gaps in the data*
- Expectation Maximization (EM) in a nutshell
  1. initialize model parameters (e.g. uniform)
  2. assign probabilities to the missing data
  3. estimate model parameters from completed data
  4. iterate steps 2–3 until convergence

# EM Algorithm


... la maison ... la maison blue ... la fleur ...



... the house ... the blue house ... the flower ...


- Initial step: all alignments equally likely
- Model learns that, e.g., **la** is often aligned with **the**

# EM Algorithm



- After one iteration
- Alignments, e.g., between **la** and **the** are more likely

# EM Algorithm



- After another iteration
- It becomes apparent that alignments, e.g., between **fleur** and **flower** are more likely (pigeon hole principle)

# EM Algorithm

... la maison ... la maison bleu ... la fleur ...  
/ | | X | |  
... the house ... the blue house ... the flower ...

- Convergence
- Inherent hidden structure revealed by EM

# EM Algorithm

21



... la maison ... la maison bleu ... la fleur ...  
... the house ... the blue house ... the flower ...

Diagram illustrating word alignment between French and English sentences. The French sentence is "... la maison ... la maison bleu ... la fleur ...". The English sentence is "... the house ... the blue house ... the flower ...". Vertical lines connect aligned words: 'la' to 'the', 'maison' to 'house', 'bleu' to 'blue', and 'fleur' to 'flower'. A large 'X' is drawn over the alignment between 'maison' and 'blue', indicating a misalignment or a correction step in the EM algorithm.

$$\begin{aligned} p(\text{la}|\text{the}) &= 0.453 \\ p(\text{le}|\text{the}) &= 0.334 \\ p(\text{maison}|\text{house}) &= 0.876 \\ p(\text{bleu}|\text{blue}) &= 0.563 \\ &\dots \end{aligned}$$

- Parameter estimation from the aligned corpus

# IBM Model 1 and EM

- EM Algorithm consists of two steps
- Expectation-Step: Apply model to the data
  - parts of the model are hidden (here: alignments)
  - using the model, assign probabilities to possible values
- Maximization-Step: Estimate model from data
  - take assign values as fact
  - collect counts (weighted by probabilities)
  - estimate model from counts
- Iterate these steps until convergence

# IBM Model 1 and EM

- We need to be able to compute:
  - Expectation-Step: probability of alignments
  - Maximization-Step: count collection

# IBM Model 1 and EM

- Probabilities

$$\begin{aligned} p(\text{the}|\text{la}) &= 0.7 & p(\text{house}|\text{la}) &= 0.05 \\ p(\text{the}|\text{maison}) &= 0.1 & p(\text{house}|\text{maison}) &= 0.8 \end{aligned}$$

- Alignments



$$p(\mathbf{e}, a|\mathbf{f}) = 0.56 \quad p(\mathbf{e}, a|\mathbf{f}) = 0.035 \quad p(\mathbf{e}, a|\mathbf{f}) = 0.08 \quad p(\mathbf{e}, a|\mathbf{f}) = 0.005$$

$$p(a|\mathbf{e}, \mathbf{f}) = 0.824 \quad p(a|\mathbf{e}, \mathbf{f}) = 0.052 \quad p(a|\mathbf{e}, \mathbf{f}) = 0.118 \quad p(a|\mathbf{e}, \mathbf{f}) = 0.007$$

- Counts

$$\begin{aligned} c(\text{the}|\text{la}) &= 0.824 + 0.052 & c(\text{house}|\text{la}) &= 0.052 + 0.007 \\ c(\text{the}|\text{maison}) &= 0.118 + 0.007 & c(\text{house}|\text{maison}) &= 0.824 + 0.118 \end{aligned}$$

# IBM Model 1 and EM: Expectation Step

- We need to compute  $p(a|\mathbf{e}, \mathbf{f})$
- Applying the chain rule:

$$p(a|\mathbf{e}, \mathbf{f}) = \frac{p(\mathbf{e}, a|\mathbf{f})}{p(\mathbf{e}|\mathbf{f})}$$

- We already have the formula for  $p(\mathbf{e}, \mathbf{a}|\mathbf{f})$  (definition of Model 1)

# IBM Model 1 and EM: Expectation Step

- We need to compute  $p(\mathbf{e}|\mathbf{f})$

$$\begin{aligned} p(\mathbf{e}|\mathbf{f}) &= \sum_a p(\mathbf{e}, a|\mathbf{f}) \\ &= \sum_{a(1)=0}^{l_f} \dots \sum_{a(l_e)=0}^{l_f} p(\mathbf{e}, a|\mathbf{f}) \\ &= \sum_{a(1)=0}^{l_f} \dots \sum_{a(l_e)=0}^{l_f} \frac{\epsilon}{(l_f + 1)^{l_e}} \prod_{j=1}^{l_e} t(e_j | f_{a(j)}) \end{aligned}$$

# IBM Model 1 and EM: Expectation Step

$$\begin{aligned} p(\mathbf{e}|\mathbf{f}) &= \sum_{a(1)=0}^{l_f} \dots \sum_{a(l_e)=0}^{l_f} \frac{\epsilon}{(l_f + 1)^{l_e}} \prod_{j=1}^{l_e} t(e_j | f_{a(j)}) \\ &= \frac{\epsilon}{(l_f + 1)^{l_e}} \sum_{a(1)=0}^{l_f} \dots \sum_{a(l_e)=0}^{l_f} \prod_{j=1}^{l_e} t(e_j | f_{a(j)}) \\ &= \frac{\epsilon}{(l_f + 1)^{l_e}} \prod_{j=1}^{l_e} \sum_{i=0}^{l_f} t(e_j | f_i) \end{aligned}$$

- Note the trick in the last line
  - removes the need for an exponential number of products
  - this makes IBM Model 1 estimation tractable

# The Trick

(case  $l_e = l_f = 2$ )

$$\begin{aligned} \sum_{a(1)=0}^2 \sum_{a(2)=0}^2 &= \frac{\epsilon}{3^2} \prod_{j=1}^2 t(e_j | f_{a(j)}) = \\ &= t(e_1 | f_0) t(e_2 | f_0) + t(e_1 | f_0) t(e_2 | f_1) + t(e_1 | f_0) t(e_2 | f_2) + \\ &\quad + t(e_1 | f_1) t(e_2 | f_0) + t(e_1 | f_1) t(e_2 | f_1) + t(e_1 | f_1) t(e_2 | f_2) + \\ &\quad + t(e_1 | f_2) t(e_2 | f_0) + t(e_1 | f_2) t(e_2 | f_1) + t(e_1 | f_2) t(e_2 | f_2) = \\ &= t(e_1 | f_0) (t(e_2 | f_0) + t(e_2 | f_1) + t(e_2 | f_2)) + \\ &\quad + t(e_1 | f_1) (t(e_2 | f_0) + t(e_2 | f_1) + t(e_2 | f_2)) + \\ &\quad + t(e_1 | f_2) (t(e_2 | f_0) + t(e_2 | f_1) + t(e_2 | f_2)) = \\ &= (t(e_1 | f_0) + t(e_1 | f_1) + t(e_1 | f_2)) (t(e_2 | f_0) + t(e_2 | f_1) + t(e_2 | f_2)) \end{aligned}$$

- Combine what we have:

$$\begin{aligned} p(\mathbf{a}|\mathbf{e}, \mathbf{f}) &= p(\mathbf{e}, \mathbf{a}|\mathbf{f})/p(\mathbf{e}|\mathbf{f}) \\ &= \frac{\frac{\epsilon}{(l_f+1)^{l_e}} \prod_{j=1}^{l_e} t(e_j|f_{a(j)})}{\frac{\epsilon}{(l_f+1)^{l_e}} \prod_{j=1}^{l_e} \sum_{i=0}^{l_f} t(e_j|f_i)} \\ &= \prod_{j=1}^{l_e} \frac{t(e_j|f_{a(j)})}{\sum_{i=0}^{l_f} t(e_j|f_i)} \end{aligned}$$

# IBM Model 1 and EM: Maximization Step

30



- Now we have to collect counts
- Evidence from a sentence pair  $\mathbf{e}, \mathbf{f}$  that word  $e$  is a translation of word  $f$ :

$$c(e|f; \mathbf{e}, \mathbf{f}) = \sum_a p(a|\mathbf{e}, \mathbf{f}) \sum_{j=1}^{l_e} \delta(e, e_j) \delta(f, f_{a(j)})$$

- With the same simplification as before:

$$c(e|f; \mathbf{e}, \mathbf{f}) = \frac{t(e|f)}{\sum_{i=0}^{l_f} t(e|f_i)} \sum_{j=1}^{l_e} \delta(e, e_j) \sum_{i=0}^{l_f} \delta(f, f_i)$$



After collecting these counts over a corpus, we can estimate the model:

$$t(e|f; \mathbf{e}, \mathbf{f}) = \frac{\sum_{(\mathbf{e}, \mathbf{f})} c(e|f; \mathbf{e}, \mathbf{f})}{\sum_e \sum_{(\mathbf{e}, \mathbf{f})} c(e|f; \mathbf{e}, \mathbf{f})}$$

# IBM Model 1 and EM: Pseudocode

**Input:** set of sentence pairs (**e**, **f**)


**Output:** translation prob.  $t(e|f)$

```
1: initialize  $t(e|f)$  uniformly
2: while not converged do
3:   // initialize
4:    $\text{count}(e|f) = 0$  for all  $e, f$ 
5:    $\text{total}(f) = 0$  for all  $f$ 
6:   for all sentence pairs (e, f) do
7:     // compute normalization
8:     for all words  $e$  in e do
9:        $\text{s-total}(e) = 0$ 
10:      for all words  $f$  in f do
11:         $\text{s-total}(e) += t(e|f)$ 
12:      end for
13:    end for
```


```
14:   // collect counts
15:   for all words  $e$  in e do
16:     for all words  $f$  in f do
17:        $\text{count}(e|f) += \frac{t(e|f)}{\text{s-total}(e)}$ 
18:        $\text{total}(f) += \frac{t(e|f)}{\text{s-total}(e)}$ 
19:     end for
20:   end for
21: end for
22: // estimate probabilities
23: for all foreign words  $f$  do
24:   for all English words  $e$  do
25:      $t(e|f) = \frac{\text{count}(e|f)}{\text{total}(f)}$ 
26:   end for
27: end for
28: end while
```

# Convergence


das Haus  
the house



das Buch  
the book



ein Buch  
a book



| $e$   | $f$  | initial | 1st it. | 2nd it. | 3rd it. | ... | final |
|-------|------|---------|---------|---------|---------|-----|-------|
| the   | das  | 0.25    | 0.5     | 0.6364  | 0.7479  | ... | 1     |
| book  | das  | 0.25    | 0.25    | 0.1818  | 0.1208  | ... | 0     |
| house | das  | 0.25    | 0.25    | 0.1818  | 0.1313  | ... | 0     |
| the   | buch | 0.25    | 0.25    | 0.1818  | 0.1208  | ... | 0     |
| book  | buch | 0.25    | 0.5     | 0.6364  | 0.7479  | ... | 1     |
| a     | buch | 0.25    | 0.25    | 0.1818  | 0.1313  | ... | 0     |
| book  | ein  | 0.25    | 0.5     | 0.4286  | 0.3466  | ... | 0     |
| a     | ein  | 0.25    | 0.5     | 0.5714  | 0.6534  | ... | 1     |
| the   | haus | 0.25    | 0.5     | 0.4286  | 0.3466  | ... | 0     |
| house | haus | 0.25    | 0.5     | 0.5714  | 0.6534  | ... | 1     |

- How well does the model fit the data?
- Perplexity: derived from probability of the training data according to the model

$$\log_2 PP = - \sum_s \log_2 p(\mathbf{e}_s | \mathbf{f}_s)$$

- Example ( $\epsilon=1$ )

|                                        | initial | 1st it. | 2nd it. | 3rd it. | ... | final  |
|----------------------------------------|---------|---------|---------|---------|-----|--------|
| $p(\text{the haus}   \text{das haus})$ | 0.0625  | 0.1875  | 0.1905  | 0.1913  | ... | 0.1875 |
| $p(\text{the book}   \text{das buch})$ | 0.0625  | 0.1406  | 0.1790  | 0.2075  | ... | 0.25   |
| $p(\text{a book}   \text{ein buch})$   | 0.0625  | 0.1875  | 0.1907  | 0.1913  | ... | 0.1875 |
| perplexity                             | 4095    | 202.3   | 153.6   | 131.6   | ... | 113.8  |

# Higher IBM Models

|             |                                |
|-------------|--------------------------------|
| IBM Model 1 | lexical translation            |
| IBM Model 2 | adds absolute reordering model |
| IBM Model 3 | adds fertility model           |
| IBM Model 4 | relative reordering model      |
| IBM Model 5 | fixes deficiency               |

- Only IBM Model 1 has global maximum
  - training of a higher IBM model builds on previous model
- Computationally biggest change in Model 3
  - trick to simplify estimation does not work anymore
  - exhaustive count collection becomes computationally too expensive
  - sampling over high probability alignments is used instead

# word alignment

# Word Alignment

Given a sentence pair, which words correspond to each other?

|         | michael | geht | davon | aus | , | dass | er | im | haus | bleibt |
|---------|---------|------|-------|-----|---|------|----|----|------|--------|
| michael |         |      |       |     |   |      |    |    |      |        |
| assumes |         |      |       |     |   |      |    |    |      |        |
| that    |         |      |       |     |   |      |    |    |      |        |
| he      |         |      |       |     |   |      |    |    |      |        |
| will    |         |      |       |     |   |      |    |    |      |        |
| stay    |         |      |       |     |   |      |    |    |      |        |
| in      |         |      |       |     |   |      |    |    |      |        |
| the     |         |      |       |     |   |      |    |    |      |        |
| house   |         |      |       |     |   |      |    |    |      |        |

# Word Alignment?

|      | john | wohnt | hier | nicht |
|------|------|-------|------|-------|
| john |      |       |      |       |
| does |      | ?     |      | ?     |
| not  |      |       |      |       |
| live |      |       |      |       |
| here |      |       |      |       |

Is the English word **does** aligned to the German **wohnt** (verb) or **nicht** (negation) or neither?

# Word Alignment?

|        | john                                                                                | biss                                                                                | ins                                                                                 | grass                                                                               |
|--------|-------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------|
| john   |  |                                                                                     |                                                                                     |                                                                                     |
| kicked |                                                                                     |  |  |  |
| the    |                                                                                     |  |  |  |
| bucket |                                                                                     |  |  |  |

How do the idioms **kicked the bucket** and **biss ins grass** match up?  
Outside this exceptional context, **bucket** is never a good translation for **grass**

# Measuring Word Alignment Quality

- Manually align corpus with *sure* ( $S$ ) and *possible* ( $P$ ) alignment points ( $S \subseteq P$ )
- Common metric for evaluation word alignments: Alignment Error Rate (AER)

$$\text{AER}(S, P; A) = 1 - \frac{|A \cap S| + |A \cap P|}{|A| + |S|}$$

- $\text{AER} = 0$ : alignment  $A$  matches all sure, any possible alignment points
- However: different applications require different precision/recall trade-offs

# symmetrization

# Word Alignment with IBM Models

- IBM Models create a **many-to-one** mapping
  - words are aligned using an alignment function
  - a function may return the same value for different input (one-to-many mapping)
  - a function can not return multiple values for one input (no many-to-one mapping)
- Real word alignments have **many-to-many** mappings

- Run IBM Model training in both directions

→ two sets of word alignment points

- Intersection: high precision alignment points
- Union: high recall alignment points
- Refinement methods explore the sets between intersection and union

## Example

english to spanish

[illegible]


spanish to english

|       |       |    |          |          |   |    |       |       |
|-------|-------|----|----------|----------|---|----|-------|-------|
|       | Maria | no | daba una | bofetada | a | la | bruja | verde |
| Mary  | ■     | □  | □        | □        | □ | □  | □     | □     |
| did   | □     | ■  | □        | □        | □ | □  | □     | □     |
| not   | □     | ■  | □        | □        | □ | □  | □     | □     |
| slap  | □     | □  | □        | ■        | □ | □  | □     | □     |
| the   | □     | □  | □        | □        | □ | ■  | □     | □     |
| green | □     | □  | □        | □        | □ | □  | □     | ■     |
| witch | □     | □  | □        | □        | □ | □  | ■     | □     |

intersection

[illegible]

# Growing Heuristics



**black:** intersection

**grey:** additional points in union

- Add alignment points from union based on heuristics:
  - directly/diagonally neighboring points
  - finally, add alignments that connect unaligned words in source and/or target
- Popular method: grow-diag-final-and