
Neural Machine Translation

Philipp Koehn

25 September 2025



Language Models




1

- Modeling variants
 - feed-forward neural network
 - recurrent neural network
 - long short term memory neural network
- May include input context

Feed Forward Neural Language Model




2



Recurrent Neural Language Model




3




Predict the first word of a sentence

Recurrent Neural Language Model




Predict the second word of a sentence
Re-use hidden state from first word prediction

Recurrent Neural Language Model



Predict the third word of a sentence
... and so on

Recurrent Neural Language Model




Recurrent Neural Translation Model



- We predicted the words of a sentence
- Why not also predict their translations?

Encoder-Decoder Model



- Obviously madness
- Proposed by Google (Sutskever et al. 2014)

What is Missing?




- Alignment of input words to output words

⇒ Solution: attention mechanism

neural translation model with attention

Input Encoding



- Inspiration: recurrent neural network language model on the input side

Hidden Language Model States


- This gives us the hidden states



- These encode left context for each word
- Same process in reverse: right context for each word




Input Encoder



- Input encoder: concatenate bidirectional RNN states
- Each word representation includes full left and right sentence context

Encoder: Math




- Input is sequence of words x_j , mapped into embedding space $\bar{E} x_j$
- Bidirectional recurrent neural networks

$$\overleftarrow{h}_j = f(\overleftarrow{h}_{j+1}, \bar{E} x_j)$$
$$\vec{h}_j = f(\vec{h}_{j-1}, \bar{E} x_j)$$

- Various choices for the function $f()$: feed-forward layer, GRU, LSTM, ...


Decoder

- We want to have a recurrent neural network predicting output words




Decoder

- We want to have a recurrent neural network predicting output words




- We feed decisions on output words back into the decoder state

- We want to have a recurrent neural network predicting output words



- We feed decisions on output words back into the decoder state
- Decoder state is also informed by the input context



- Decoder is also recurrent neural network over sequence of hidden states s_i

$$s_i = f(s_{i-1}, E y_{i-1}, c_i)$$


- Again, various choices for the function $f()$: feed-forward layer, GRU, LSTM, ...
- Output word y_i is selected by computing a vector t_i (same size as vocabulary)

$$t_i = W(U s_{i-1} + V E y_{i-1} + C c_i)$$

then finding the highest value in vector t_i


- If we normalize t_i , we can view it as a probability distribution over words
- $E y_i$ is the embedding of the output word y_i

Attention




- Given what we have generated so far (decoder hidden state)
- ... which words in the input should we pay attention to (encoder states)?

Attention



- Given: – the previous hidden state of the decoder s_{i-1}
– the representation of input words $h_j = (\overleftarrow{h}_j, \overrightarrow{h}_j)$
- Predict an alignment probability $a(s_{i-1}, h_j)$ to each input word j (modeled with with a feed-forward neural network layer)


Attention



- Normalize attention (softmax)


$$\alpha_{ij} = \frac{\exp(a(s_{i-1}, h_j))}{\sum_k \exp(a(s_{i-1}, h_k))}$$

Attention



- Relevant input context: weigh input words according to attention: $c_i = \sum_j \alpha_{ij} h_j$


Attention



- Use context to predict next hidden state and output word

training


Comparing Prediction to Correct Word




- Current model gives some probability $t_i[y_i]$ to correct word y_i
- We turn this into an error by computing cross-entropy: $-\log t_i[y_i]$

Computation Graph

- Math behind neural machine translation defines a computation graph
- Forward and backward computation to compute gradients for model training




Unrolled Computation Graph



- Already large degree of parallelism
 - most computations on vectors, matrices
 - efficient implementations for CPU and GPU
- Further parallelism by batching
 - processing several sentence pairs at once
 - scalar operation $\rightarrow$ vector operation
 - vector operation $\rightarrow$ matrix operation
 - matrix operation $\rightarrow$ 3d tensor operation
- Typical batch sizes 50–100 sentence pairs

Batches


- Sentences have different length
- When batching, fill up unneeded cells in tensors



⇒ A lot of wasted computations

Mini-Batches

- Sort sentences by length, break up into mini-batches



- Example: Maxi-batch 1600 sentence pairs, mini-batch 80 sentence pairs


Overall Organization of Training

- Shuffle corpus
- Break into maxi-batches
- Break up each maxi-batch into mini-batches
- Process mini-batch, update parameters
- Once done, repeat
- Typically 5-15 epochs needed (passes through entire training corpus)

deeper models

Deeper Models


- Encoder and decoder are recurrent neural networks
- We can add additional layers for each step
- Recall shallow and deep language models



- Adding residual connections (short-cuts through deep layers) help

Deep Decoder

- Two ways of adding layers
 - deep transitions: several layers on path to output
 - deeply stacking recurrent neural networks
- Why not both?



Deep Encoder

- Previously proposed encoder already has 2 layers
 - left-to-right recurrent network, to encode left context
 - right-to-left recurrent network, to encode right context

⇒ Third way of adding layers

